

# Rethinking Security from Scratch

The Case for Shifting Container  
Security from the Edge to the Core

Chris Milsted

Staff Field Engineer

[cmilsted@vmware.com](mailto:cmilsted@vmware.com)

+44 7795 316 723

# Agenda

\$ whoami

Refresher - OCI image specification

Options - Distroless and Scratch

Demo - Building from scratch

Looking at our current security model

Shifting towards “build-based” security

\$ whoami

# \$ whoami

- Chris is based in the UK and is a Staff Field Engineer for VMware working in the Cloud Native team. He spends most of his work wrangling Kubernetes and most of his spare time playing field hockey badly and being a taxi driver for two children who are growing up rapidly.
- He has spent the last 5 years working with containers and kubernetes and Linux
- Most of this has been with large Financial Services Customers

# What inspired this talk

Use the smallest base image possible

The base image is the one referenced in the `FROM` instruction in your Dockerfile. Every other instruction in the Dockerfile builds on top of this image. The smaller the base image, the smaller the resulting image is, and the more quickly it can be downloaded. For example, the `alpine:3.7` image is 71 MB smaller than the `centos:7` image.

You can even use the `scratch` base image, which is an empty image on which you can build your own runtime environment. If your app is a statically linked binary, it's easy to use the scratch base image:

```
FROM scratch
COPY mybinary /mybinary
CMD [ "/mybinary" ]
```



zwischenzugs

## A Docker Image in Less Than 1000 Bytes

5 Minutes

vmware®



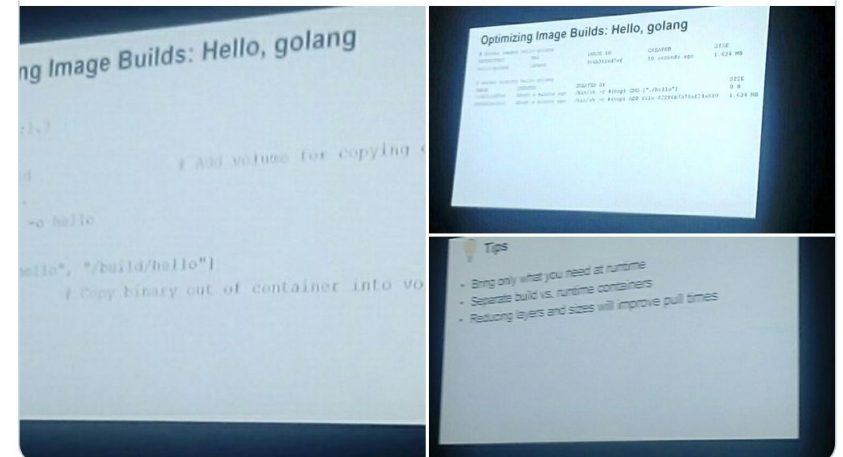
**Kelsey Hightower** @kelseyhightower · 1 Dec 2016

This is great advice for keeping **container images small**.



**AWS re:Invent** @AWSreInvent · 1 Dec 2016

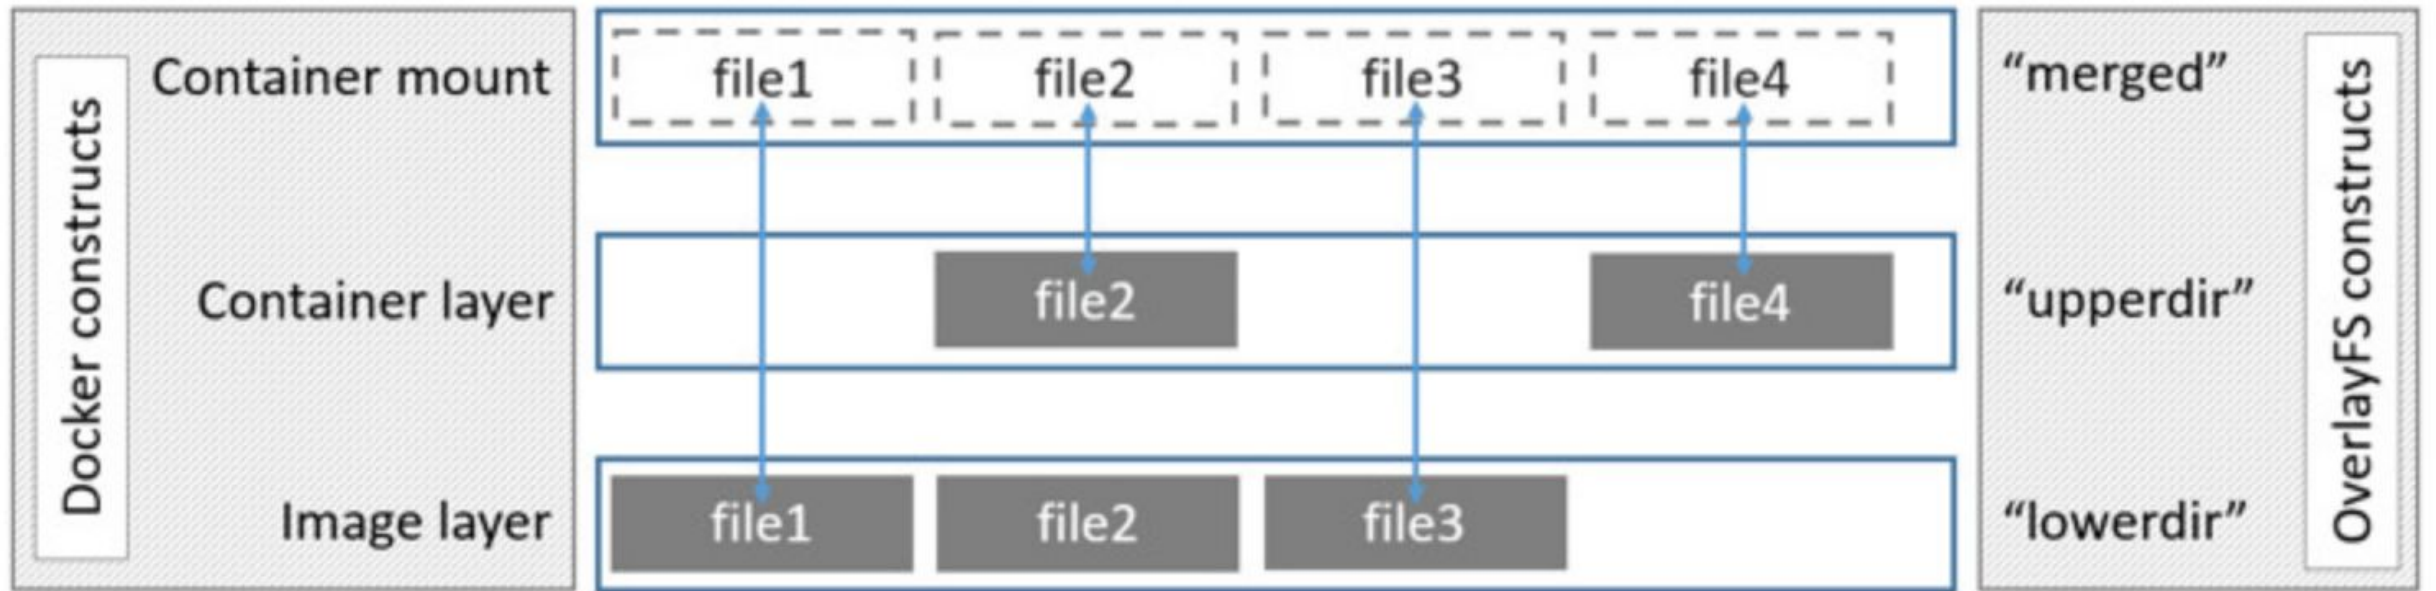
Using a GO language example, at tip noted for optimization is to separate your Docker build vs. runtime containers #reInvent #CON401



# Refresher - OCI image specification

# What is a container

## Overlay2 graph driver



# OCI images

## Building container images and inheritance

alpine:latest

golang:alpine

alpine:latest

```
$ sudo docker inspect --format='{{.GraphDriver.Data}}' alpine:latest |tr ' ' '\n'  
map[MergedDir:/var/lib/docker/overlay2/784af3f8492d8d7ade0a82bbaa6dace2bd694d4c0f1a4ab1510cd43cec0c67d9/merged  
UpperDir:/var/lib/docker/overlay2/784af3f8492d8d7ade0a82bbaa6dace2bd694d4c0f1a4ab1510cd43cec0c67d9/diff  
WorkDir:/var/lib/docker/overlay2/784af3f8492d8d7ade0a82bbaa6dace2bd694d4c0f1a4ab1510cd43cec0c67d9/work]
```

```
$ sudo docker inspect --format='{{.GraphDriver.Data}}' golang:alpine |tr ' ' '\n'  
map[LowerDir:/var/lib/docker/overlay2/c34762b4720a23034e53c9f35be009071921086ca140db65d13a82940b9ebf35/diff:/var/lib/docker/overlay2/9988592682bef4c04e144f6d153b9116686e8c0a879e71e41cfc79e07037a19d/diff:/var/lib/docker/overlay2/96f7522bc8c562a266b524645dd43931d6bdf37560f940fe50cb85177f08fe02/diff:/var/lib/docker/overlay2/784af3f8492d8d7ade0a82bbaa6dace2bd694d4c0f1a4ab1510cd43cec0c67d9/diff  
MergedDir:/var/lib/docker/overlay2/71495101f982ea3ee928f5c00798c7ee6b6ae21378d40f252e9283b856c792d6/merged  
UpperDir:/var/lib/docker/overlay2/71495101f982ea3ee928f5c00798c7ee6b6ae21378d40f252e9283b856c792d6/diff  
WorkDir:/var/lib/docker/overlay2/71495101f982ea3ee928f5c00798c7ee6b6ae21378d40f252e9283b856c792d6/work]
```



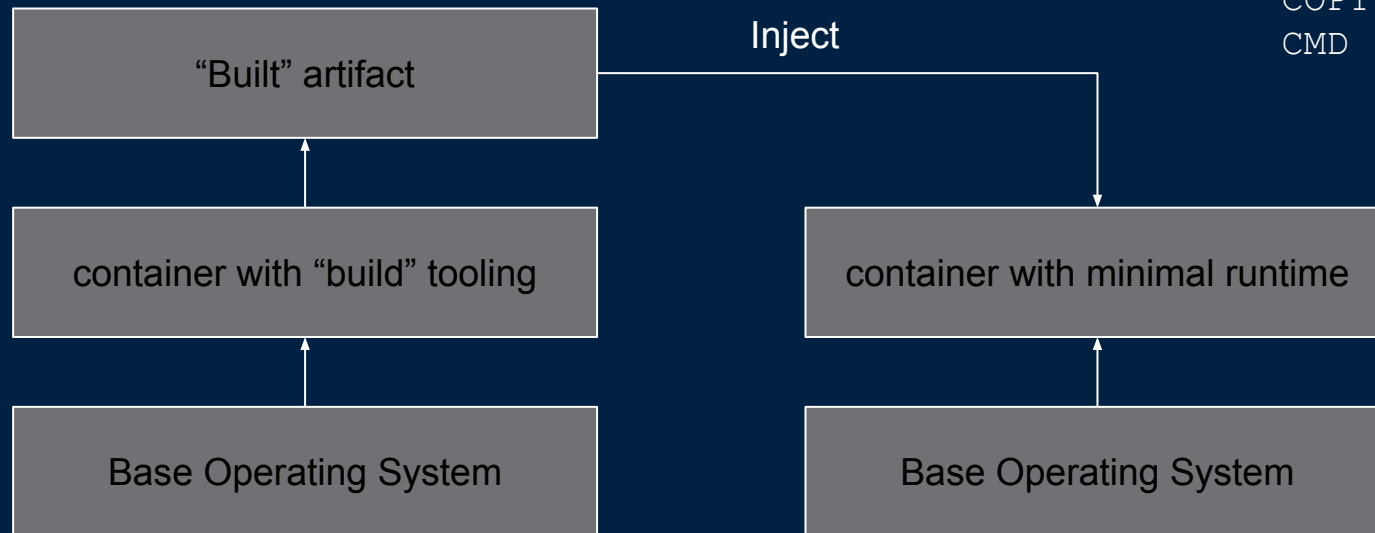
# Options - Distroless and Scratch

# Distroless


Minimising container but still based on distro userspace

```
FROM golang:1.7.3 AS builder
WORKDIR /go/src/github.com/<pathto code>
COPY app.go .
RUN CGO_ENABLED=0 GOOS=linux go build -a
-installsuffix cgo -o app .
```

```
FROM alpine:latest
RUN apk --no-cache add ca-certificates
WORKDIR /root/
COPY --from=builder ./app .
CMD ["/app"]
```



# Scratch



```
FROM scratch
```

```
# List the maintainer  
LABEL maintainer="Chris Milsted"
```

```
# Copy the Pre-built binary file from the  
previous stage.  
COPY ./helloworld .
```

```
# Expose port 8080 to the outside world  
EXPOSE 8080
```

```
#Command to run the executable  
CMD ["./helloworld"]
```

# Layer Isolation

An application fully isolated in a container must be self-contained

- Namespaces
- CGroups
- Kernel capabilities (seccomp profile)
- Linux Security Module (e.g. SELinux or apparmor)

# Demo - Building from scratch

ubuntu@master-1:~/go/src/helloworld\$ ls -l

```
total 16
-rw-rw-r-- 1 ubuntu ubuntu 292 Dec 30 12:40 Dockerfile
-rw-rw-r-- 1 ubuntu ubuntu 292 Nov 28 10:45 Dockerfile.orig
-rw-rw-r-- 1 ubuntu ubuntu 306 Dec 30 12:39 Dockerfile.static
-rw-rw-r-- 1 ubuntu ubuntu 374 Nov 27 05:07 helloworld.go
```

ubuntu@master-1:~/go/src/helloworld\$ cat helloworld.go

```
package main
```

```
import (
    "fmt"
    "net/http"
)
```

```
func helloHandler(w http.ResponseWriter, r *http.Request) {
    fmt.Fprintln(w, "Hello World from Go in minimal Docker container")
}
```

```
func main() {
    http.HandleFunc("/", helloHandler)

    fmt.Println("Started, serving at 8080")
    err := http.ListenAndServe(":8080", nil)
    if err != nil {
        panic("ListenAndServe: " + err.Error())
    }
}
ubuntu@master-1:~/go/src/helloworld$
```

```
ubuntu@master-1:~/go/src/helloworld$ ldd helloworld
linux-vdso.so.1 (0x00007ffcecff4000)
libpthread.so.0 => /lib/x86_64-linux-gnu/libpthread.so.0 (0x00007f39aa77f000)
libc.so.6 => /lib/x86_64-linux-gnu/libc.so.6 (0x00007f39aa38e000)
/lib64/ld-linux-x86-64.so.2 (0x00007f39aa99e000)
```

Typically, compiled **Go** binaries have very few dependencies. You do not need any kind of runtime libraries or VMs, and all **Go** libraries that you use in your project are embedded directly into the resulting executable file. However, if you compile your application in Linux, the **Go** compiler will link the resulting binary against a few C standard libraries that are typically available on any Linux system. If you are on Linux, you can easily find out against which libraries your program was linked by invoking the `ldd` binary with one of your compiled **Go** binaries as argument. If your binary is linked against the C standard library, you will receive the following output:

```
$ ldd ./eventservice
linux-vdso.so.1 (0x00007ffed09b1000)
libpthread.so.0 => /lib/x86_64-linux-gnu/libpthread.so.0
(0x00007fd523c36000)
libc.so.6 => /lib/x86_64-linux-gnu/libc.so.6 (0x00007fd52388b000)
/lib64/ld-linux-x86-64.so.2 (0x0000564d70338000)
```

This means that your **Go** application actually needs these Linux libraries to run and that you cannot just arbitrarily delete them from your image to make it smaller.





ubuntu@master-1: ~/go/src/helloworld



ubuntu@master-1: ~/go/src/helloworld

chris@chris-Precision-5530: ~/Videos

```
ubuntu@master-1:~/go/src/helloworld$ ls -lrt
total 7292
-rw-rw-r-- 1 ubuntu ubuntu    374 Nov 27 05:07 helloworld.go
-rw-rw-r-- 1 ubuntu ubuntu    292 Nov 28 10:45 Dockerfile.orig
-rw-rw-r-- 1 ubuntu ubuntu    306 Dec 30 12:39 Dockerfile.static
-rw-rw-r-- 1 ubuntu ubuntu    292 Dec 30 12:40 Dockerfile
-rwxrwxr-x 1 ubuntu ubuntu 7450480 Dec 30 14:11 helloworld
ubuntu@master-1:~/go/src/helloworld$ env CGO_ENABLED=0 GOOS=linux GOARCH=amd64 go build -a -o helloworld_static
ubuntu@master-1:~/go/src/helloworld$ vim Dockerfile
ubuntu@master-1:~/go/src/helloworld$ sudo docker build .
Sending build context to Docker daemon 14.85MB
Step 1/5 : FROM scratch
-->
Step 2/5 : LABEL maintainer="Chris Milsted"
--> Using cache
--> 7880abf6e518
Step 3/5 : COPY ./helloworld_static .
--> 5245b122de8b
Step 4/5 : EXPOSE 8080
--> Running in 8f5555007efb
Removing intermediate container 8f5555007efb
--> f8509693a0d4
Step 5/5 : CMD ["/helloworld_static"]
--> Running in 0ea9c07e3794
Removing intermediate container 0ea9c07e3794
--> fba8d2dc11e1
Successfully built fba8d2dc11e1
ubuntu@master-1:~/go/src/helloworld$ sudo docker run --rm -p 8080:8080 fba8d2dc11e1
Started, serving at 8080
```



# Success?

A very very minimal container - which just contains a single executable

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
<none>	<none>	fba8d2dc11e1	3 minutes ago	7.39MB
<none>	<none>	92769f32e1f8	41 minutes ago	7.45MB

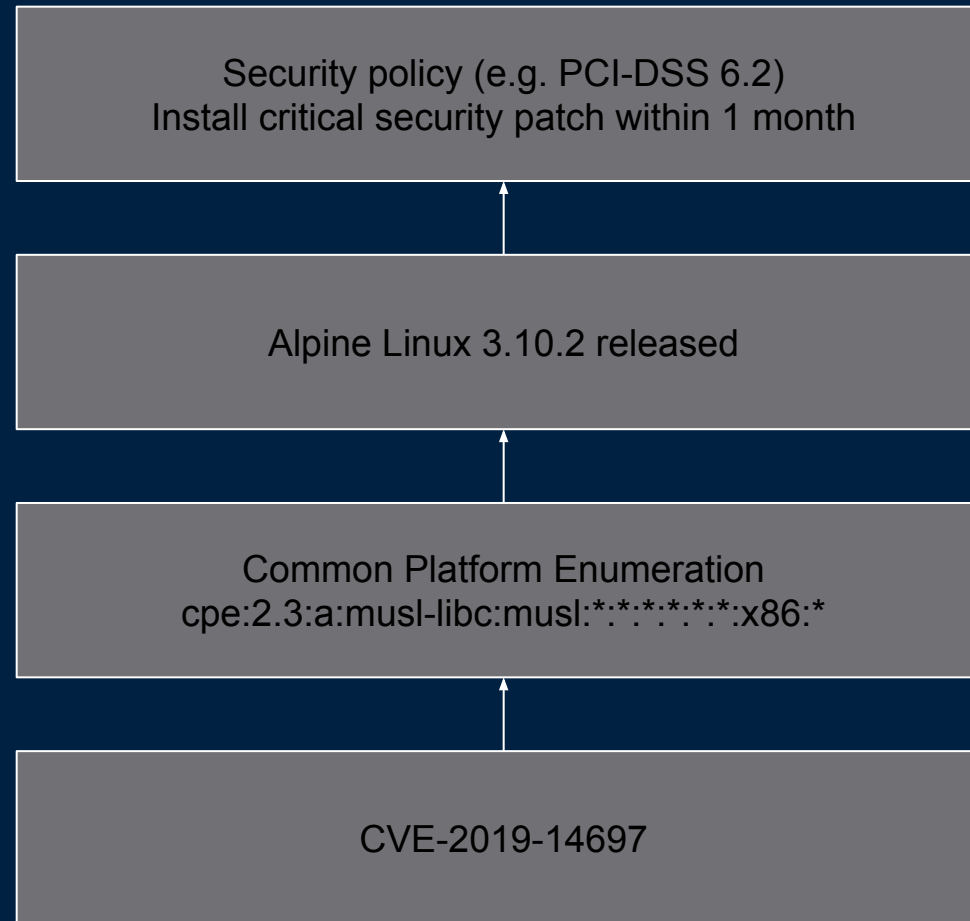
but...

```
ubuntu@master-1:~/go/src/helloworld$ sudo docker run --rm -p 8080:8080 fba8d2dc11e1 sh
docker: Error response from daemon: OCI runtime create failed: container_linux.go:346: starting container process caused "exec: \"sh\": executable file not found in $PATH": unknown.
ERRO[0001] error waiting for container: context canceled
```

What is inside our container when we scale this approach...

```
(gdb) p 'runtime.buildVersion'
$1 = 0x700d3a "go1.13.4"
```

# Looking at our current security model



Process maps CVE and code change to package version and then onto higher level component such as Operating System.

# Shifting towards “build-based” security

# A quest for minimal build (technology) has impacted a process

Today we noticed that another Go team at CloudFlare, the [Data](#) team, had **a much smarter way to bake version numbers into binaries using the -X linker option.**

The -X [Go linker option](#), which you can set with `-ldflags`, sets the value of a string variable in the Go program being linked. You use it like this: `-X main.version 1.0.0`.

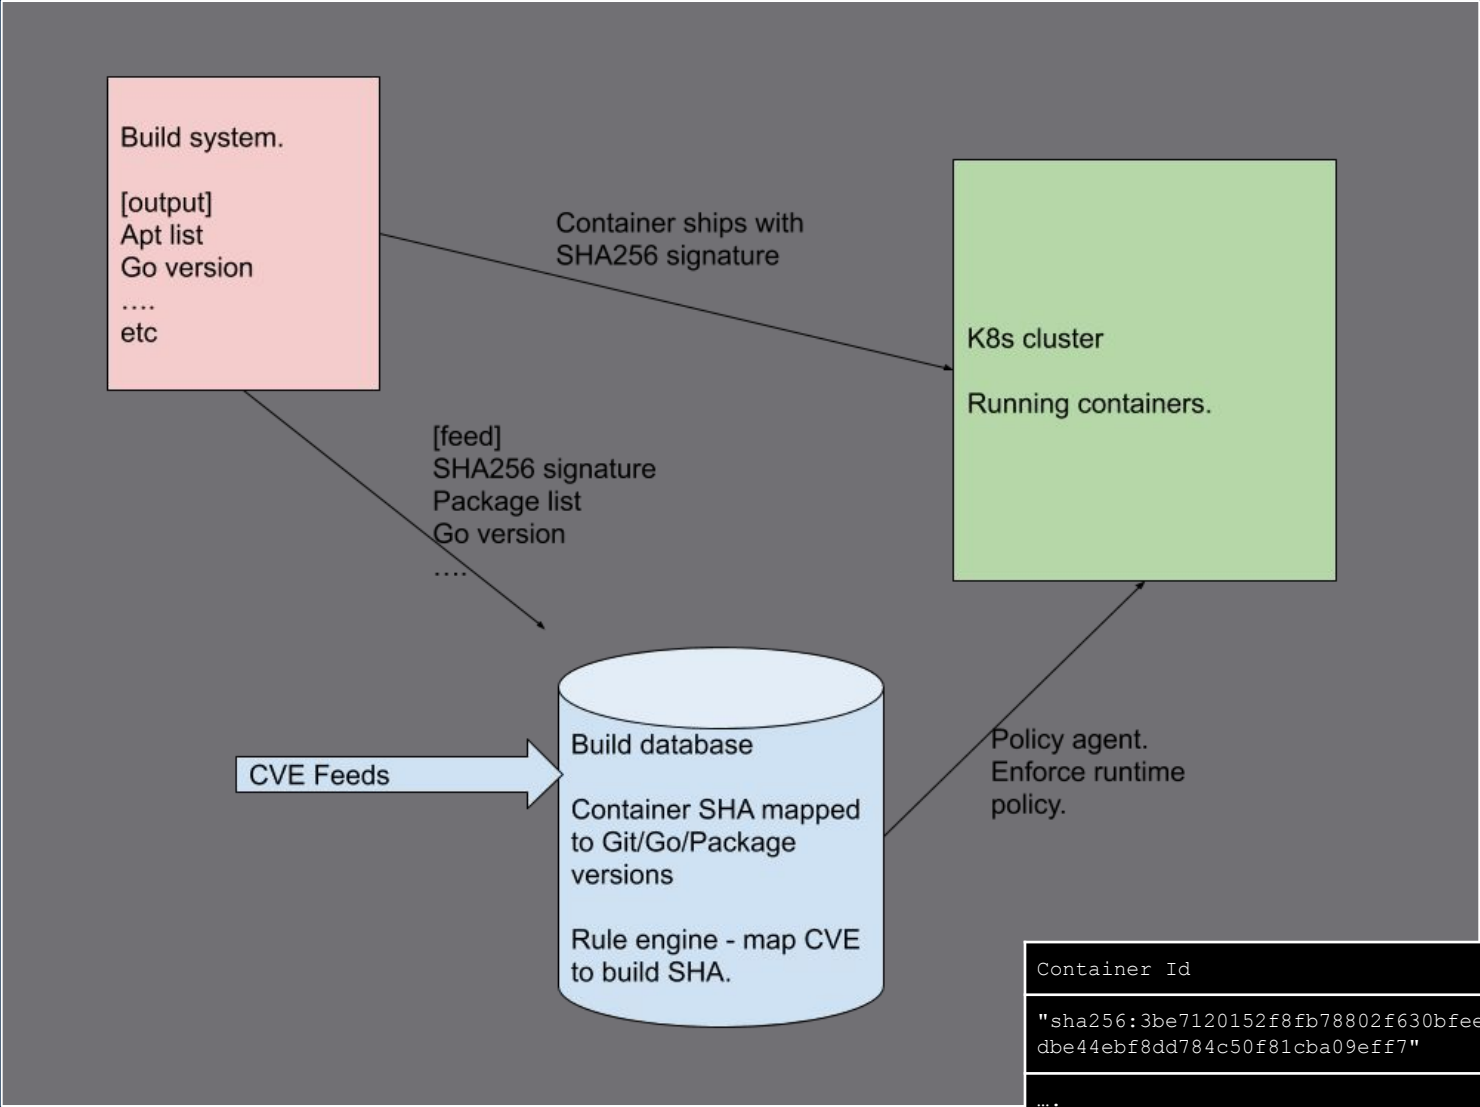
## Content-based non-linear versioning

The very fundamental foundation of everything at Polyverse is content-based versioning. *A [content-based version](#) is a cryptographically secure hash over a set of bits, and that hash is used as the "version" for those bits.*

This is a premise you will find everywhere in our processes. If you want to tell your colleague what version of Router you want, you'd say something like: `router@72e5e550d1835013832f64597cb1368b7155bd53`. That is the version of the router you're addressing. It is unambiguous, and you can go to the Git repository that holds our router, and get PRECISELY what your colleague is using

# What does the future hold?

We trust the build system, not the edge container



Container Id	Go	Git tag	libc
"sha256:3be7120152f8fb78802f630bfeec68af821dbe44ebf8dd784c50f81cba09eff7"	go1.13.4 linux/amd64	6a615ca	2.27-3ubuntu1
...			

# Questions?



# Thank You